

CS 5435:

Security and Privacy Concepts in the Wild

Homework #1

Due: Before class on 17 Sept. 2015

90 points

+ 20 bonus points (applied to term bonus point pool, not to homework grade)

Instructor: Ari Juels TA: Fan Zhang

Problem 1: Name three incentives an attacker might have for paying \$325 (or more) for a high-profile Twitter account password. At least one of these should make money for the attacker. [15 points]

Problem 2: Recall that many password-reset questions are weak, in the sense of having a high guessing probability.

Question 2.1: For an online service whose users are U.S. inhabitants, give an estimate of the average number of randomly selected accounts an attacker would have to try to answer the question “What’s your mother’s maiden name?” correctly. Assume one guess per account. (Hint: Use statistics available at http://www.census.gov/topics/population/genealogy/data/2000_surnames.html or <http://1.usa.gov/1EKssO1>.) [5 points]

Question 2.2: Find the weakest password-recovery question you can online that wasn’t discussed in class. Give an estimate of the guessing probability and explain how you arrived at it. [12 points]

Problem 3: The Ruritania National Bank (RNB) encrypts the four-digit PINs of its customers under a fixed, secret, randomly generated 256-bit key K . It may (or may not) be helpful to know how the RNB encrypts a PIN P for user U :

The bank converts the PIN into a big-endian binary representation, pads it out by appending 0s to yield a 128-bit plaintext block, and then encrypts this block under key K using AES-256. The notorious Zenda League hacked the RNB to protest the bank's refusal to provide free checking accounts to its customers. They posted a spreadsheet with the encrypted PINs of 10,000 of RNB's customers, represented in hexadecimal. You'll find this spreadsheet (in the file RNBencryptedPINs.xlsx) on the course site.

Question 3.1: Suppose an attacker that knows the RNB's message formatting and encryption scheme is given a single encrypted PIN randomly selected from the spreadsheet. How many decryption attempts, expressed as a decimal number, would the attacker need to perform to learn the PIN with probability 1? (You may assume heuristically that the attacker knows when decryption has been successful; it can with overwhelming probability by, e.g., testing on other spreadsheet values.) [4 points]

Question 3.2: Now, to show a different attack that efficiently obtains information about customer PINs, answer the following question. What PIN belongs to the customer Ebenezer Scrooge? Note that a '*' in the customer name column denotes a customer other than Ebenezer Scrooge. (Hint: See <http://www.datagenetics.com/blog/september32012/> or <http://bit.ly/1cZJIvL>.) [10 points]

Problem 4: The NYC Taxi Authority recently released records of taxi trips in which medallion numbers were concealed via hashing with MD5. (See the *Ars Technica* article at <http://bit.ly/1o0196n>.) You can find a full list of current medallion numbers at <https://data.cityofnewyork.us/Transportation/Current-Medallions/avwq-z233> or <http://bit.ly/1ITn06z>. (So can a potential attacker. . .)

Question 4.1: What is the maximum number of attempts needed to crack the hash of a medallion number (assuming no precomputation and no database of issued medallions)? (Hint: The number of possible medallion numbers reported in the *Ars Technica* article is wrong. Another hint: There are three authorized medallion formats, DLDD, LLDDD, and LLLDDD, where D stands for 'digit' and L for 'letter'.) [5 points]

Question 4.2: What is the security goal of the NYC Department of Transportation in this setting, i.e., why did they hash medallion numbers? [3 points]

Question 4.3: On the course site (file medallionhashes.txt), you'll find a list of SHA-256 hashes for some LLLDDD-format medallions that are active at the time of this exercise. Crack them! Also indicate who owns the corresponding taxis. (Hint: To verify the correctness of the SHA-256 code you use, here's the hash of medallion SBV381:

95fa2dcf62a7ad26b29a08a03fe16c7a3aa341afb4e62999e12b6311d0d43f6a.

The SHA-256 option at www.quickhash.com may also be helpful.) [15 points]

Question 4.4: Suppose that VIN numbers and licensee names were concatenated with medallion numbers prior to hashing active medallion numbers. Would this fix the problem and achieve the desired security goal? Why or why not? [5 points]

Question 4.5: Design a cryptographic scheme (a mechanism) using SHA-256 that would permit NYC to generate a unique pseudonym for each taxi medallion and meet the desired security goal. (Note: Given NYC's difficulties, maybe there's a startup or at least an interesting project somewhere in here. . .) [5 points]

Problem 5: Oklahoma Otis, the famous gambler, wants to start a gambling site for single-player roulette. (A standard American wheel has 38 slots. Otis thinks there's a business for an innovative, bigger wheel, one that instead has 39 slots.) He designed the following pseudocode client-server cryptographic protocol for his site.

$C \rightarrow S$:	cslot	% client submits bet cslot
S :	$y = \text{truerand}()$;	
	$\text{sslot} = y \bmod 39$;	% server chooses random slot
	if cslot = sslot, then	
	win = true	% client wins
	else	
	win = false	% client loses
$S \rightarrow C$:	sslot, win	% send winning slot and outcome to client

(Some notation: C stands for client, S for server. The pseudocode $X \rightarrow Y$ denotes transmission of a message from entity X to Y , while $X:$ denotes a local computation by entity X . The function SHA256 implements SHA-256; you may assume it acts as a random oracle. The function $\text{truerand}()$ returns a 256-bit random string and the symbol \parallel , used below, denotes string concatenation.)

Question 5.1: Otis's friend, Mississippi Mabel, the famous sharpshooter, shot down his idea. She said his protocol won't work because the server can cheat without being detected. What is the security goal of the client? [2 points] How can the server defeat this security goal? [2 points] Is there any way in Otis's protocol for the client to detect server cheating? [2 points]

Question 5.2: Mabel came up with a clever protocol to enable the client to detect a cheating server immediately. Here is her protocol. (The 256-bit key K was generated at random.)

```

K =      7c49 70fe 78bd 92ea e9c2 22db f690 5ffe
        35ab 2c65 2a4d c6d4 bc52 9880 e133 8be2

C :      a = SHA256(cslot  $\parallel$  K)
C  $\rightarrow$  S : a
S :      y = truerand();
        sslot = y mod 39
S  $\rightarrow$  C : sslot
C  $\rightarrow$  S : cslot
S :      if a = SHA256(cslot  $\parallel$  K) AND cslot = sslot then
            win = true
        else
            win = false
S  $\rightarrow$  C : win
C :      if cslot = sslot AND win = false
            cheated = true
        else
            cheated = false

```

Why doesn't Mabel's protocol work, i.e., how can a server ensure that it never awards a win to a client but $\text{cheated} = \text{false}$? [5 points]

Question 5.3 (Bonus): Fix Mabel’s protocol so that it does that does work, i.e., it allows a client to detect unfair play immediately. [10 points]

Problem 6 (Bonus): Recall that a MAC is a “tag” that may be appended to a message (plaintext or ciphertext) to ensure that tampering can be detected by the receiver. MACs are typically constructed using hash functions. A common example is HMAC.

The Chief Cryptographer of the RNB used SHA-256 to design a new, custom MAC called RuriMAC. RuriMAC has two good security properties: (1) It cannot be forged (with certainty) by an eavesdropper without intercepting a complete, valid (message, MAC) pair, and (2) Even after seeing many valid (message, MAC) pairs, an adversary cannot feasibly compute the secret key. Unfortunately, though an adversary that intercepts a single valid (message, MAC) pair can forge another valid (message, MAC) pair. Give a possible design for RuriMAC. (Yes, we are asking you to construct a *bad* MAC scheme!) [10 points]